SchemaX-070: Food Donations Matching System Database Design

Data Definition Language Commands

```
CREATE TABLE User (
    user_id BIGINT AUTO_INCREMENT PRIMARY KEY,
    name VARCHAR(1024) NOT NULL,
    location VARCHAR(2048),
    contact_preference VARCHAR(255),
    email VARCHAR(255) NOT NULL,
    phone VARCHAR(255) NOT NULL,
    type VARCHAR(128)
);

CREATE TABLE Donor (
    donor_id BIGINT AUTO_INCREMENT PRIMARY KEY,
    type VARCHAR(255),
    preferred_pickup_time VARCHAR(255),
    user_id BIGINT NOT NULL,
    FOREIGN KEY (user_id) REFERENCES User(user_id)
);

CREATE TABLE Recipient (
    recipient_id BIGINT AUTO_INCREMENT PRIMARY KEY,
    preferences VARCHAR(2048),
    user_id BIGINT NOT NULL,
    notification_enabled BOOLEAN,
    FOREIGN KEY (user_id) REFERENCES User(user_id)
);

CREATE TABLE Item (
    item_id BIGINT AUTO_INCREMENT PRIMARY KEY,
    item_name VARCHAR(255),
    category VARCHAR(255)
);

CREATE TABLE Listing (
    listing_id BIGINT AUTO_INCREMENT PRIMARY KEY,
    listed_by BIGINT NOT NULL,
    location VARCHAR(2048),
    expiration_date DATE,
```

```
    type VARCHAR(255),
    pickup_time_range VARCHAR(255),
    status VARCHAR(255),
    FOREIGN KEY (listed_by) REFERENCES Donor(donor_id)
);

CREATE TABLE Booking (
    booking_id BIGINT AUTO_INCREMENT PRIMARY KEY,
    booked_by BIGINT NOT NULL,
    booking_status VARCHAR(255) NOT NULL,
    pickup_datetime Date(255),
    FOREIGN KEY (booked_by) REFERENCES Recipient(recipient_id)
);

CREATE TABLE ListingItem (
    listing_item_id BIGINT AUTO_INCREMENT PRIMARY KEY,
    quantity BIGINT,
    item_id BIGINT NOT NULL,
    listing_id BIGINT NOT NULL,
    booking_id BIGINT,
    expiration_date DATE,
    status VARCHAR(255),
    FOREIGN KEY (item_id) REFERENCES Item(item_id),
    FOREIGN KEY (listing_id) REFERENCES Listing(listing_id),
    FOREIGN KEY (booking_id) REFERENCES Booking(booking_id)
);

CREATE TABLE Review (
    review_id BIGINT AUTO_INCREMENT PRIMARY KEY,
    review TEXT,
    review_date DATE,
    rating INT,
    donor_id BIGINT NOT NULL,
    recipient_id BIGINT NOT NULL,
    FOREIGN KEY (donor_id) REFERENCES Donor(donor_id),
    FOREIGN KEY (recipient_id) REFERENCES Recipient(recipient_id)
);

CREATE TABLE Recommendation (
    matching_id BIGINT AUTO_INCREMENT PRIMARY KEY,
    type VARCHAR(255),
    food_type VARCHAR(255),
    recipient_id BIGINT NOT NULL,
    donor_id BIGINT NOT NULL,
```

```
    score DOUBLE,
    FOREIGN KEY (recipient_id) REFERENCES Recipient(recipient_id),
    FOREIGN KEY (donor_id) REFERENCES Donor(donor_id)
);
```

Database Hosting (GCP), Setup and Connection

The database server is hosted on GCP with minimum configurations as suggested in the workshop lecture.

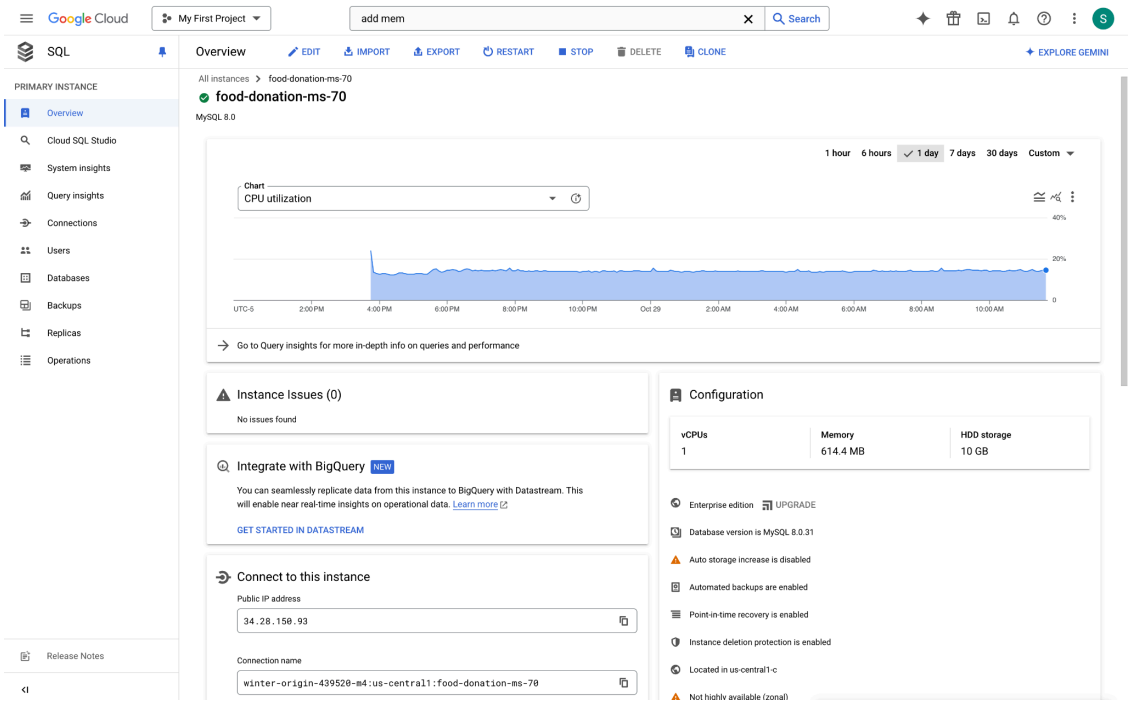

Here’s the connection through GCP Cloud Shell

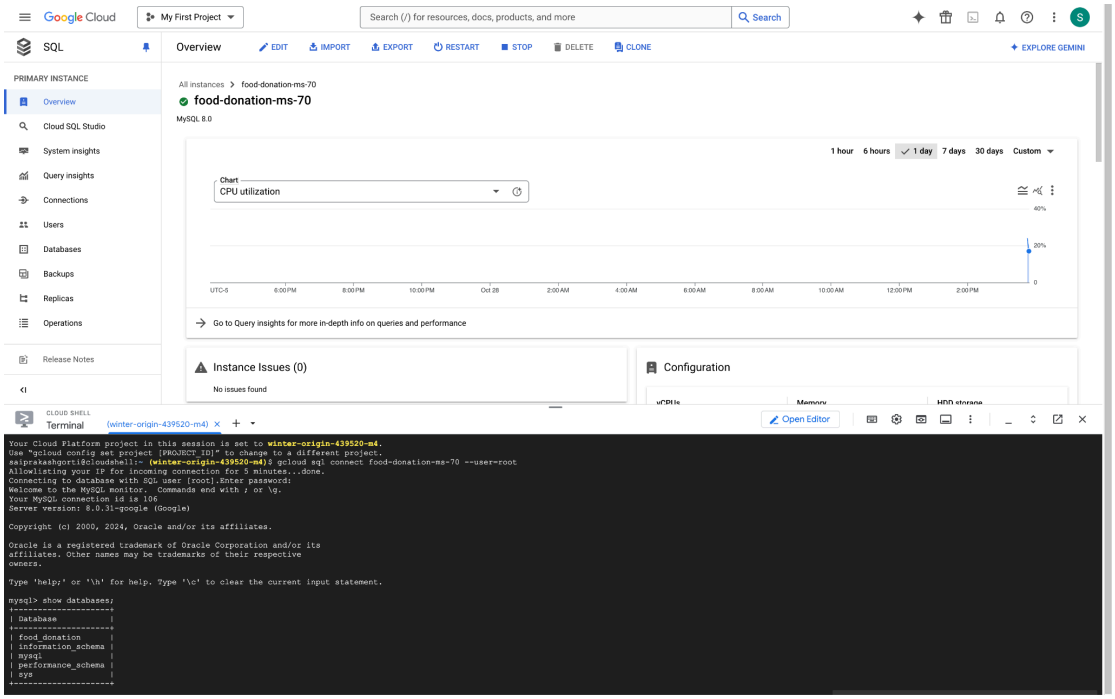

Moving forward, we’ll use Datagrip to connect to this server for a better user interface when writing queries. Here’s the connection to the MySQL server hosted on GCP through Datagrip.

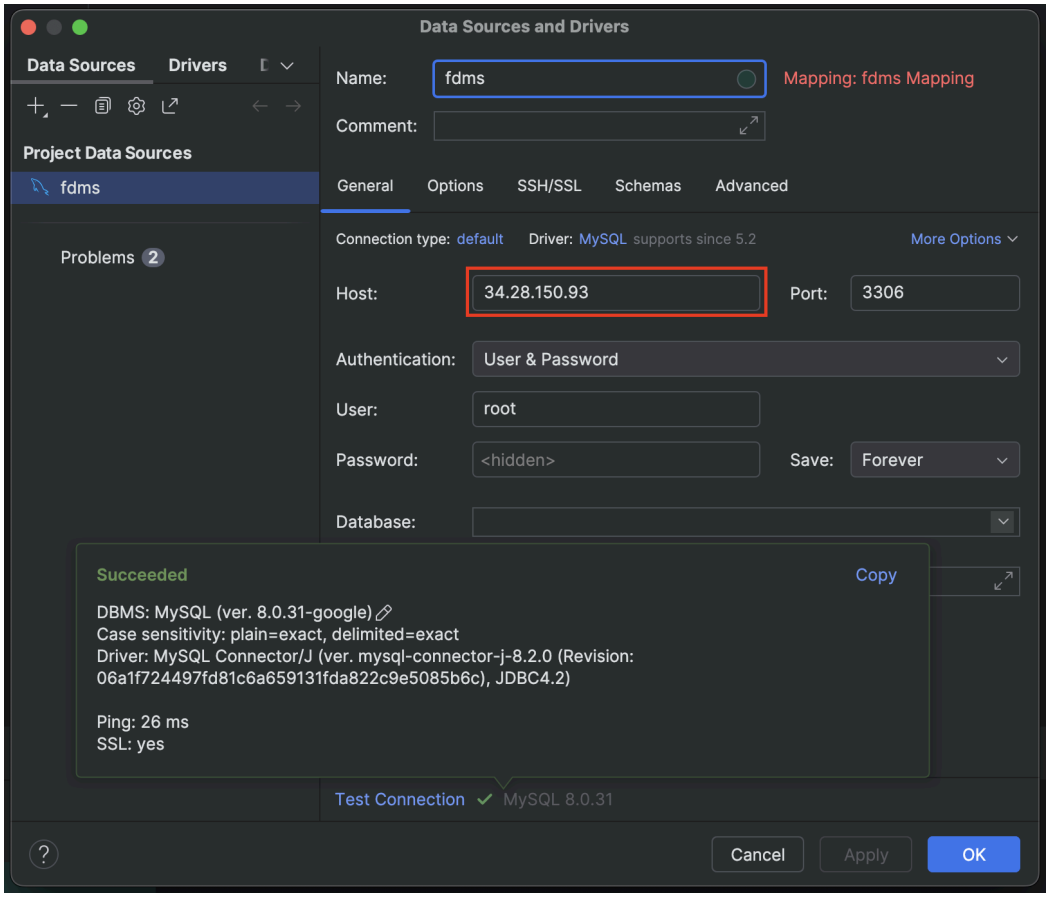

Data Insertion

The data inserted into the tables is a combination of the datasets mentioned in the previous stage and synthetic data to get optimal results. Here's the number of rows we inserted into each table.

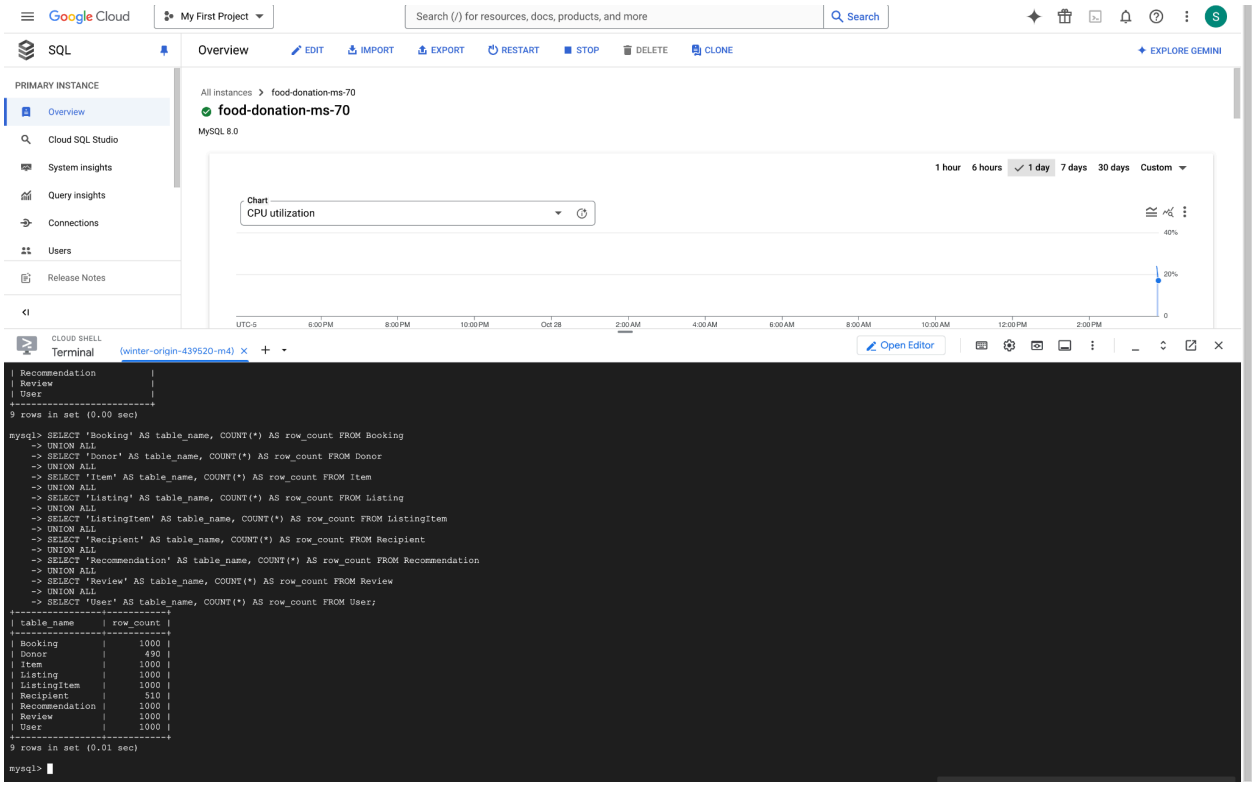

Advanced SQL Queries

AQ1: Listing Status Summary for Donor

This query is to provide a summary of all the listings and their associated items for a specific donor. The key purpose is to keep track of the status of all listings and the number of items within each listing.

Note: The query will have 2 Complex SQL relations: JOINS, AGGREGATION

```
SELECT l.listing_id,
       li.status,
       COUNT(*) AS numberOfItems,
       SUM(Quantity) AS totalQuantity
FROM Listing l
LEFT JOIN ListingItem li ON l.listing_id = li.listing_id
```

```
WHERE l.listed_by = '<donor_id>'
GROUP BY l.listing_id, status;
```

Below is the same query run on actual data.

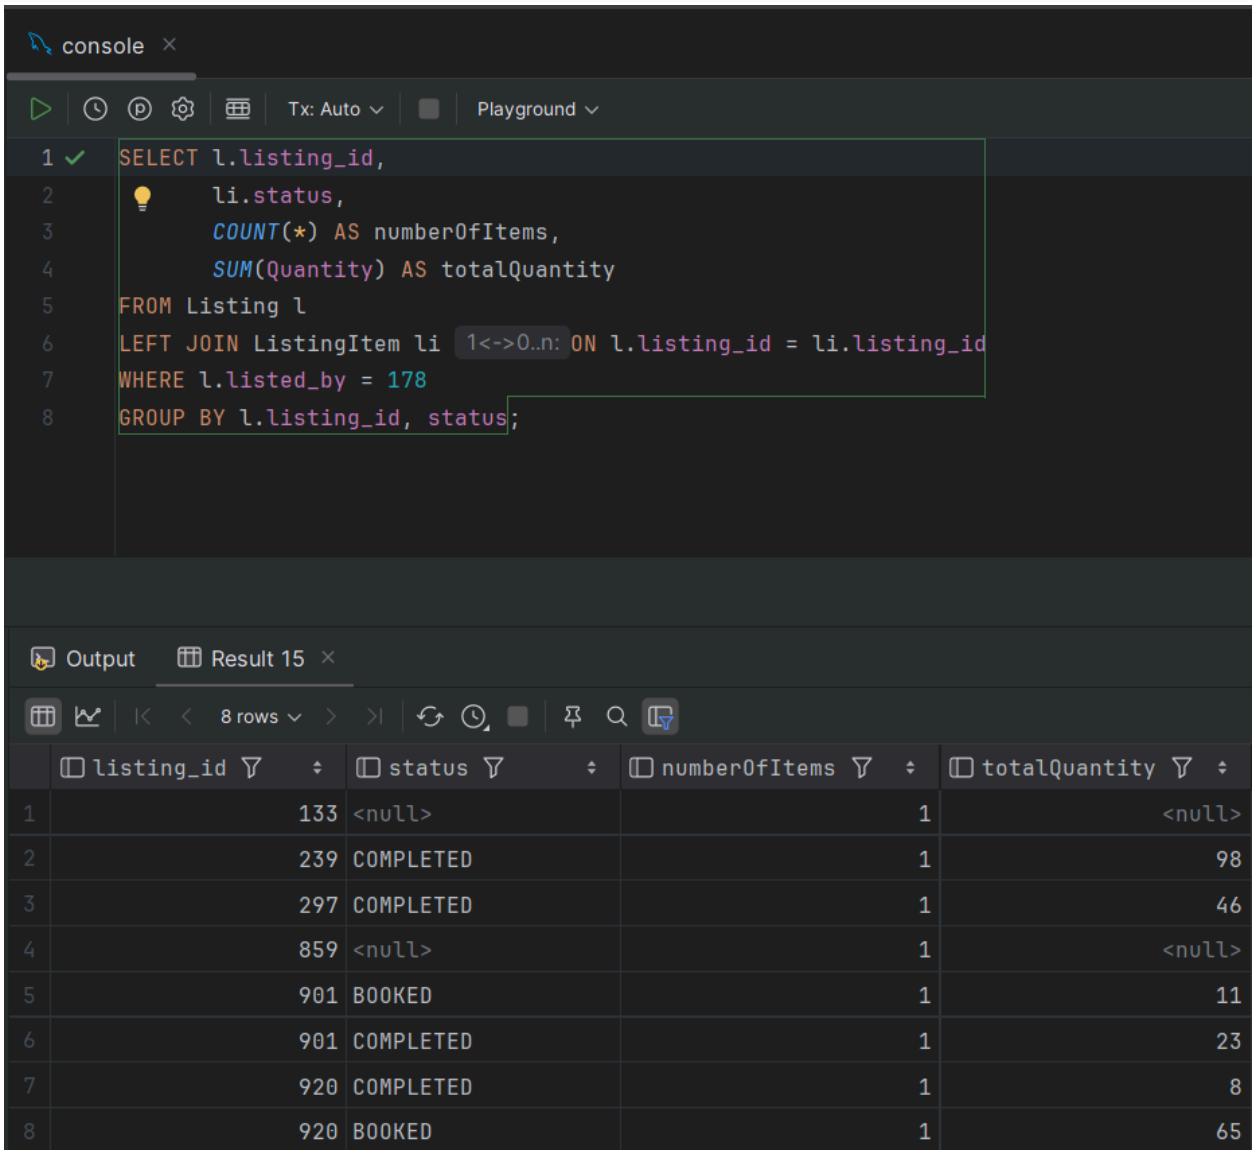

AQ2: Search-based Listing Recommendations

This query combines search functionality with a recommendations table to fetch relevant listings having the best recommendation scores by using nested queries to filter out the items based on search terms (ordered by best score) and then joining listings with recommendations on donor_id.

Note: The query will have 2 Complex SQL relations: JOINS, NESTED SUBQUERY

```sql
SELECT DISTINCT l.listed_by,
                l.listing_id,
                l.location,
                l.status
FROM Listing l
LEFT JOIN (
    SELECT donor_id, score
    FROM Recommendation
    WHERE recipient_id = <RECIPIENT_ID>
) AS rec ON l.listed_by = rec.donor_id
WHERE listing_id IN (
    SELECT DISTINCT listing_id
    FROM ListingItem li
    WHERE li.item_id IN (
        SELECT DISTINCT item_id
        FROM Item
        WHERE item_name = '<ITEM_SEARCH_QUERY>'
            OR category = '<CATEGORY_SEARCH_QUERY>'
    )
    ORDER BY score DESC
);
```

Below is the same query run on actual data. The resulting output contains more than 15 rows; hence, LIMIT was used to show 15 rows.

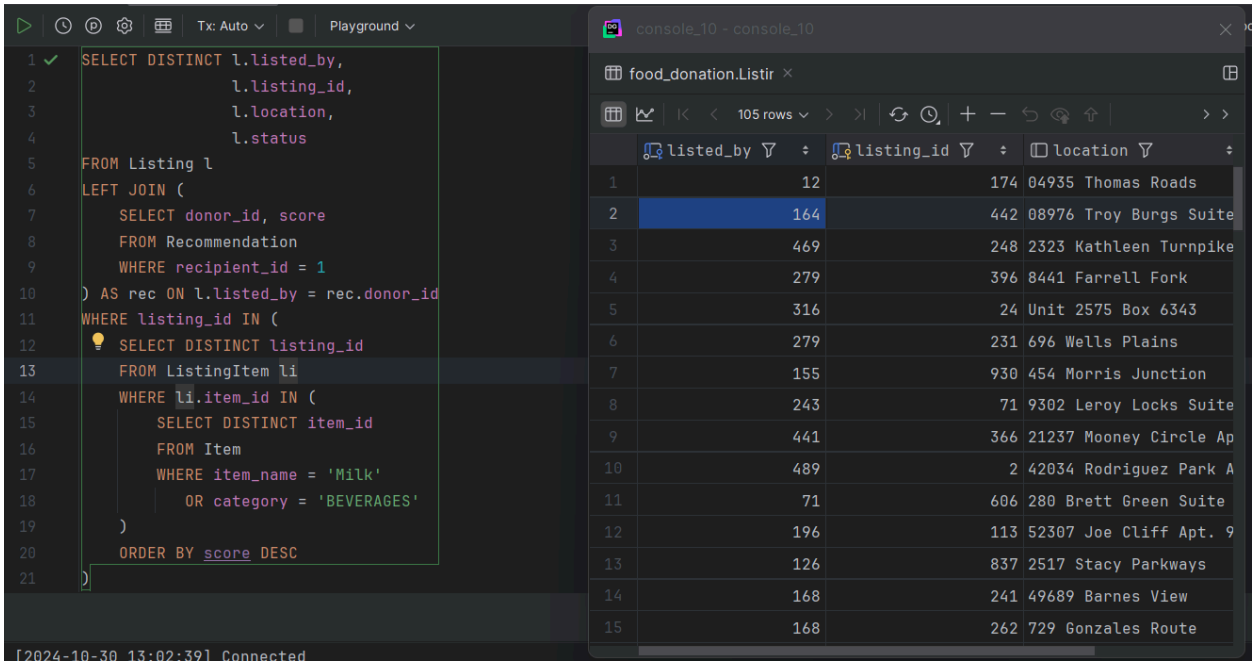

AQ3: Fetch Detailed Donor Profile

The below query fetches the donors details, along with the average rating and overall number of reviews. This will be used on the donor detail view page, where we fetch the user details along with the review summary. We can also add filter to filter for a specific donor_id using where clause as per need.

Note: The query will have 3 Complex SQL relations: JOINS, AGGREGATIONS, SUBQUERY

```sql
SELECT d.donor_id,
       u.name, u.location, u.email, u.phone,
       rev.avg_rating,
       rev.review_count
FROM Donor d
LEFT JOIN User u ON d.user_id = u.user_id
LEFT JOIN (
   SELECT donor_id,
          Avg(rating) AS avg_rating,
          Count(*) AS review_count
   FROM Review r
   GROUP BY r.donor_id
) AS rev
ON d.donor_id = rev.donor_id
WHERE avg_rating IS NOT NULL AND review_count IS NOT NULL
```

Below is the same query run on actual data. The resulting output contains more than 15 rows; hence, LIMIT was used to show 15 rows.

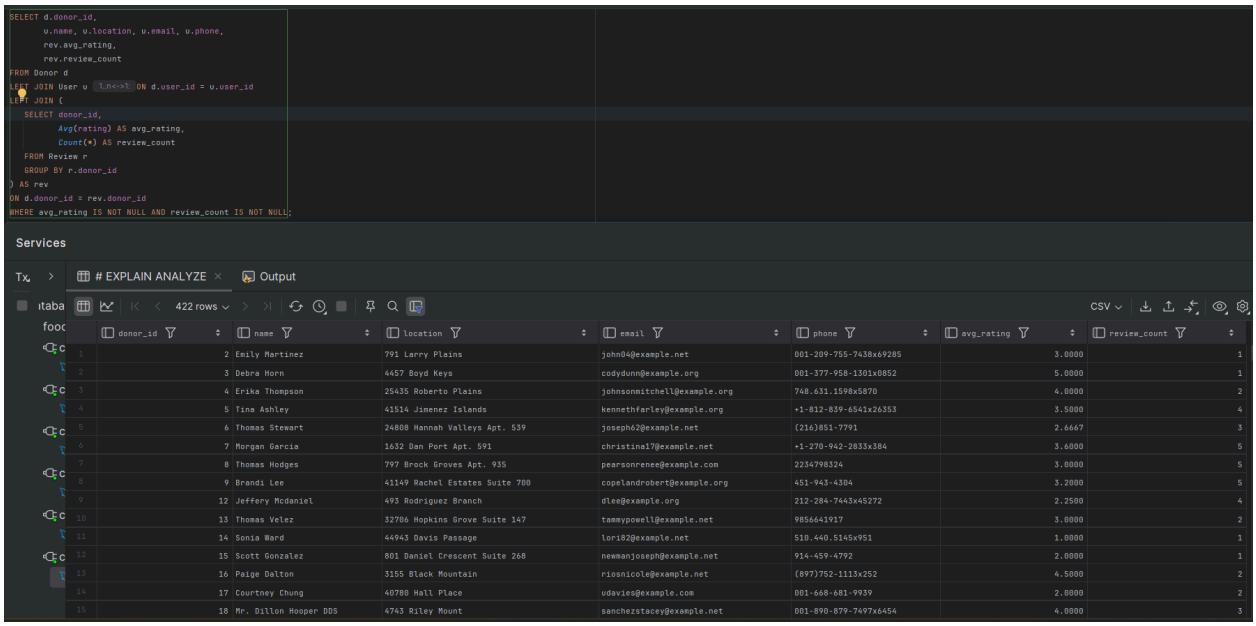

	donor_id	name	location	email	phone	avg_rating	review_count
1	2	Emily Martinez	791 Larry Plains	john04@example.net	001-209-755-7438x69285	3.0000	1
2	3	Debra Horn	4457 Boyd Keys	codydunn@example.org	001-377-958-1301x0852	5.0000	1
3	4	Erika Thompson	25435 Roberto Plains	johnsonmitchell@example.org	748.631.1598x5870	4.0000	2
4	5	Tina Ashley	41514 Jimenez Islands	kennethfarley@example.org	+1-812-839-6541x26353	3.5000	4
5	6	Thomas Stewart	24808 Hannah Valleys Apt. 539	joseph62@example.net	(216)851-7791	2.6667	3
6	7	Morgan Garcia	1632 Dan Port Apt. 591	christina17@example.net	+1-270-942-2833x384	3.6000	5
7	8	Thomas Hodges	797 Brock Groves Apt. 935	pearsonrenee@example.com	2234798324	3.0000	5
8	9	Brandi Lee	41149 Rachel Estates Suite 700	copelandrobert@example.org	451-943-4304	3.2000	5
9	12	Jeffery Mcdaniel	493 Rodriguez Branch	dlee@example.org	212-284-7443x45272	2.2500	4
10	13	Thomas Velez	32706 Hopkins Grove Suite 147	tammypowell@example.net	9856641917	3.0000	2
11	14	Sonia Ward	44943 Davis Passage	lori82@example.net	510.440.5145x951	1.0000	1
12	15	Scott Gonzalez	801 Daniel Crescent Suite 268	newmanjoseph@example.net	914-459-4792	2.0000	1
13	16	Paige Dalton	3155 Black Mountain	riosnicole@example.net	(897)752-1113x252	4.5000	2
14	17	Courtney Chung	40780 Hall Place	udavies@example.com	001-668-681-9939	2.0000	2
15	18	Mr. Dillon Hooper DDS	4743 Riley Mount	sanchezstacey@example.net	001-890-879-7497x6454	4.0000	3

AQ4: Fetch Available Listings by Recipient Preferences

This query fetches all the available listings based on the preferences set by users. The main idea is to split all the comma-separated preferences and then search for these in the listings items table. Also, need to ensure that either the product has no expiry date, or it has not expired yet. Another check would be to check if each listing is ACTIVE or not, and if the listing_item is AVAILABLE or not.

Note: This query requires multiple JOIN, NESTED SUBQUERY.

```
SELECT l.listing_id, l.location,
       li.item_id,
       i.item_name,
       i.category,
       li.quantity, li.expiration_date,
       d.donor_id,
       u.name AS donor_name
FROM Listing l
JOIN ListingItem li ON l.listing_id = li.listing_id
JOIN Item i ON li.item_id = i.item_id
JOIN Donor d ON l.listed_by = d.donor_id
JOIN User u ON d.user_id = u.user_id
JOIN Recipient r ON r.recipient_id = <RECIPIENT_ID>
WHERE l.status = 'ACTIVE' AND li.status = 'AVAILABLE'
  AND (li.expiration_date IS NULL OR li.expiration_date > CURRENT_DATE)
  AND CONCAT(',', r.preferences, ',') LIKE CONCAT('%,', i.category, ',%')
```

We can extend this by adding a sort-by option and appending the following to the query

```
ORDER BY li.expiration_date ASC,
         l.pickup_time_range ASC;
```

Below is the same query run on actual data. The resulting output contains more than 15 rows; hence, LIMIT was used to show 15 rows.

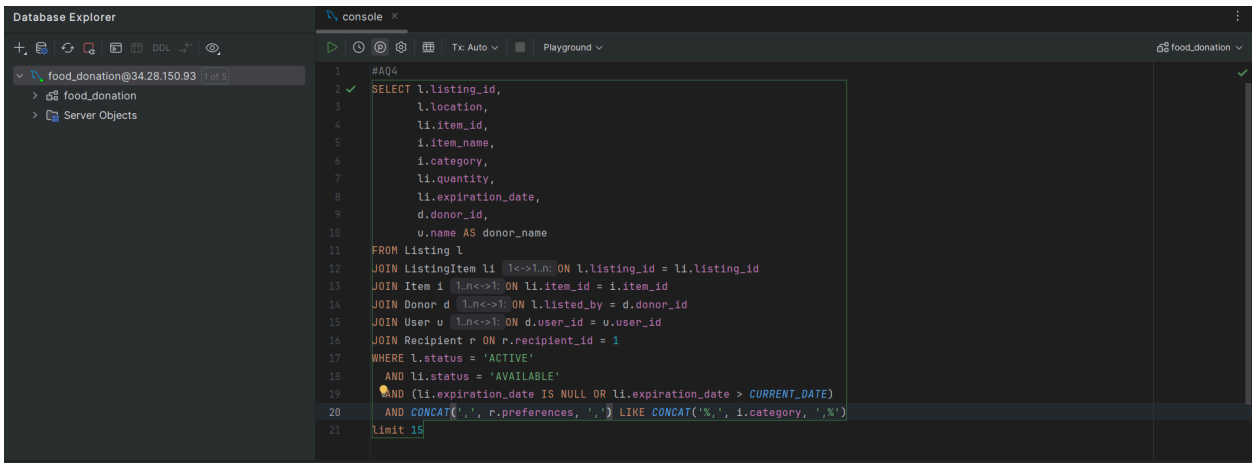

Output #AQ4

15 rows CSV

	listing_id	location	item_id	item_name	category	quantity	expiration_date	don
1	290	6334 Christina Haven Apt. 975	523	Value Pack Pasta Sauce - Bulk	DAIRY	81	2024-11-26	
2	259	045 Michael Island Suite 548	809	Regular Whole Milk - Medium	DAIRY	94	2024-11-25	
3	571	8984 Nicole Orchard Apt. 393	5	Regular Whole Milk - Small	BABY FOOD	64	2024-12-13	
4	2	42034 Rodriguez Park Apt. 721	791	Value Pack Chicken Breast - Medium	BABY FOOD	8	2024-11-28	
5	581	312 Downs Bypass Suite 104	389	Regular Chicken Breast - Medium	DAIRY	93	2024-12-31	
6	895	333 Rebecca Ridge Apt. 992	264	Organic Fresh Apples - Small	DAIRY	78	2025-01-16	
7	717	105 Jasmine Mount	904	Organic Pasta Sauce - Small	DAIRY	13	2024-11-18	
8	789	28596 Charles Plains Suite 155	570	Premium Baby Formula - Medium	DAIRY	59	2024-12-07	
9	514	0635 Ramirez Crescent	254	Premium Orange Juice - Large	DAIRY	56	2024-12-02	
10	768	500 Donald Knoll Apt. 193	285	Regular Fresh Bananas - Small	DAIRY	64	2025-01-15	
11	514	0635 Ramirez Crescent	723	Value Pack Carrots - Bulk	BABY FOOD	35	2025-01-16	
12	636	53835 Rodriguez Greens	595	Value Pack Instant Oatmeal - Large	DAIRY	89	2025-01-18	
13	470	7645 Bill Views Apt. 910	533	Regular Peanut Butter - Large	DAIRY	17	2024-11-26	
14	43	38475 Cassandra Vista Suite 195	571	Value Pack Baby Formula - Medium	DAIRY	85	2024-11-02	
15	908	531 Brown Bridge	650	Premium Orange Juice - Medium	BABY FOOD	97	2025-01-20	

Indexing: Design and Analysis

AQ1: Index Design and Analysis

The EXPLAIN ANALYZE command BEFORE Indexing is shown below:

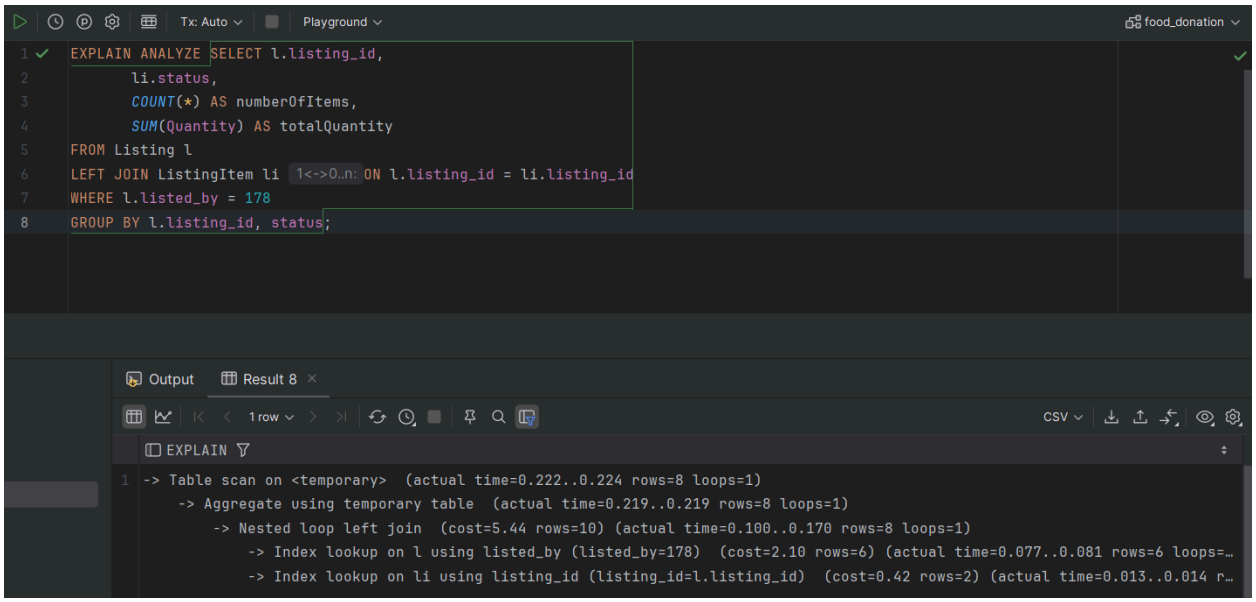

Let's explore some Index designs for the AQ1:

1) Creating index on ListingItem.Status: No effect was observed on the overall cost.

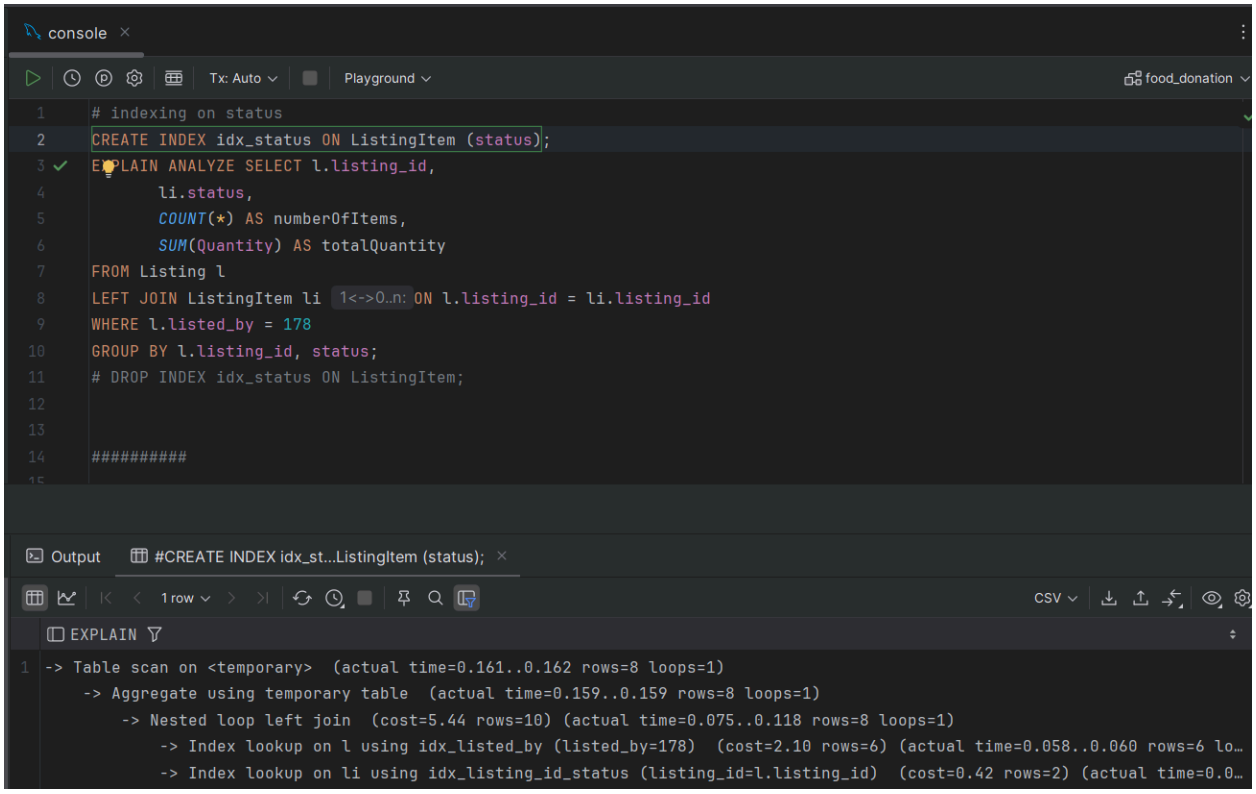

2) Creating composite index on listing_id and Status: No effect observed on overall cost.

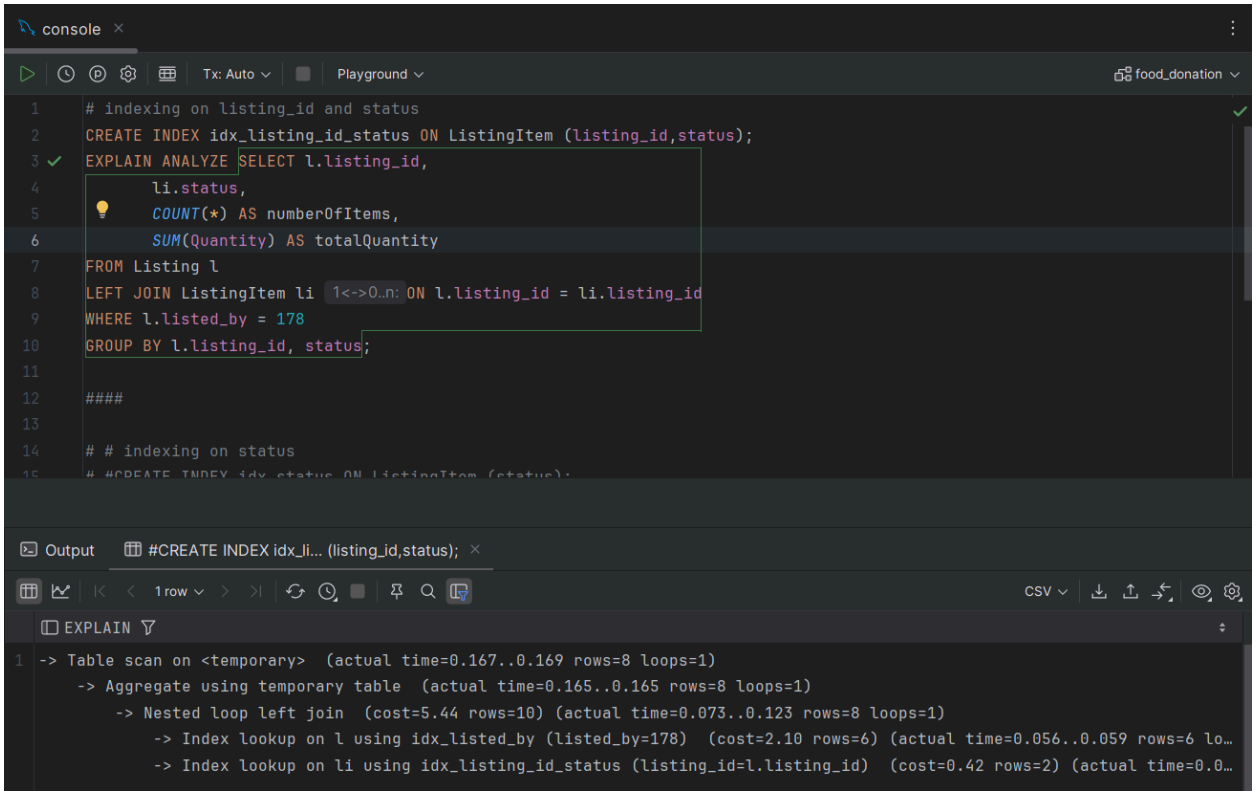

3) Creating Index on Quantity: The overall cost remains contant even with Quantity index

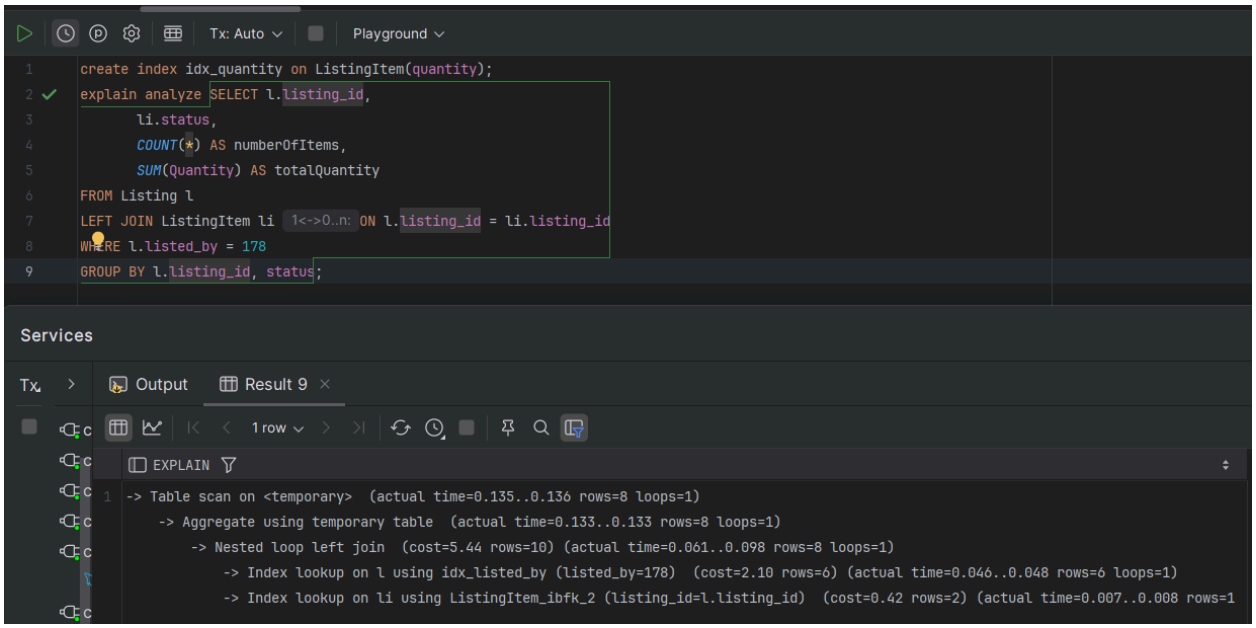

Analysis/Reason for no change in cost:
Despite creating a composite index on **listing_id** and **status**, no improvement was observed in the query's performance. This can be attributed to the query structure, which utilizes a LEFT JOIN and a filter condition (WHERE l.listed_by = '<donor_id>') that does not involve the indexed fields. Therefore, the index on listing_id and status is not fully utilized because the query's

filtering condition is based on **listed_by** (**already a foreign key index**), making the composite index less relevant for optimizing this specific query. Additionally, aggregating and grouping operations (**COUNT** and **SUM**) may already be efficiently handled by the database engine without significant dependence on the specified index, further reducing the index's impact. Also, the cardinality of the **status** column is less, which might be the reason behind no change since we are anyway fetching that many rows. Also, indexing on quantity did not bring that much impact. The reason might be that it is a high cardinality column which is used inside aggregation. As we can see, the index was not even used in the query plan.

Note: The above query **only** has a **listing_id** and **status** field on which the Indexing can be applied on. Other columns are already foreign or primary key, hence we can't index on those..

AQ2: Index Design and Analysis

For AQ2, the EXPLAIN ANALYZE command BEFORE Indexing is shown below:

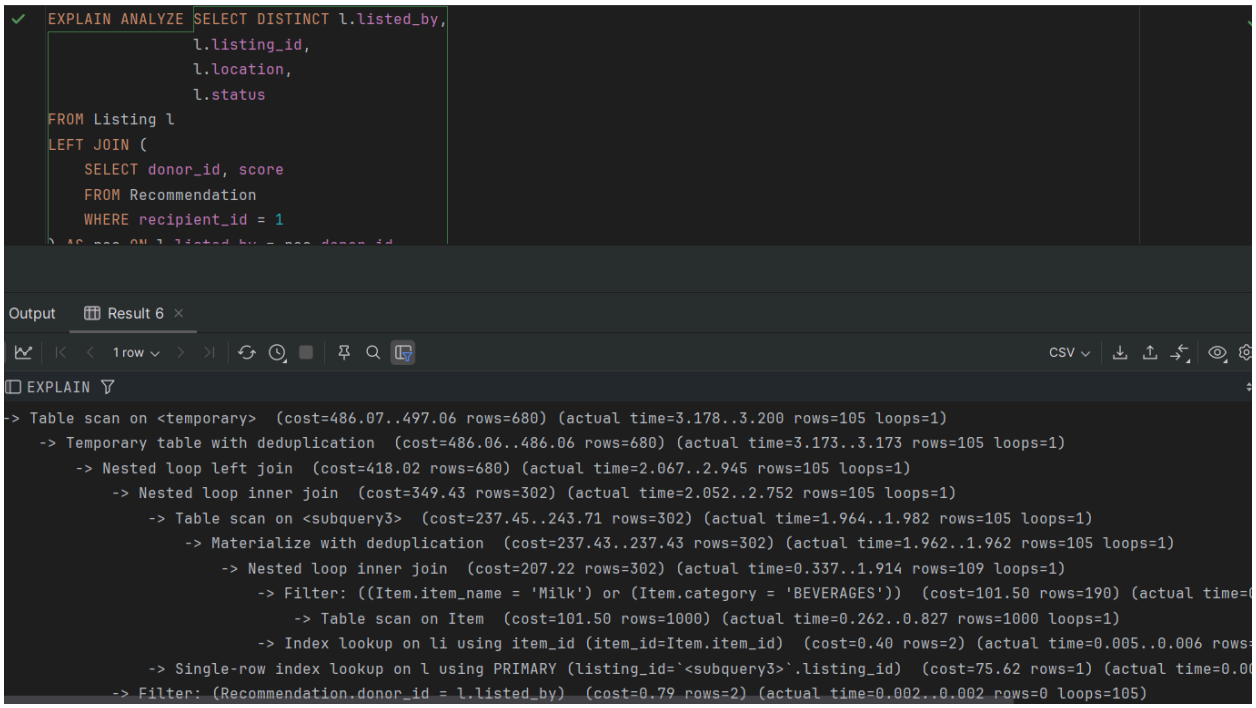

Let's explore the index designs for AQ2:

1) Indexing on Item.item_name only:

```
CREATE INDEX idx_item_name_only ON Item(item_name);

EXPLAIN ANALYZE SELECT DISTINCT l.listed_by,
                l.listing_id,
                l.location,
                l.status
FROM Listing l
LEFT JOIN (
    SELECT donor_id, score
```

Output | Result 10-2 | # drop index idx_ite...ame_category on Item;

EXPLAIN

```
-> Table scan on <temporary>  (cost=309.85..316.90 rows=366) (actual time=1.562..1.576 rows=105 loops=1)
    -> Temporary table with deduplication  (cost=309.83..309.83 rows=366) (actual time=1.560..1.560 rows=105 loops=1)
        -> Nested loop left join  (cost=273.22 rows=366) (actual time=1.063..1.429 rows=105 loops=1)
            -> Nested loop inner join  (cost=236.05 rows=163) (actual time=1.054..1.301 rows=105 loops=1)
                -> Table scan on <subquery3>  (cost=174.67..179.16 rows=163) (actual time=1.041..1.078 rows=105 loops=1)
                    -> Materialize with deduplication  (cost=174.64..174.64 rows=163) (actual time=1.040..1.040 rows=105 loops=1)
                        -> Nested loop inner join  (cost=158.39 rows=163) (actual time=0.101..1.009 rows=109 loops=1)
                            -> Filter: ((Item.item_name = 'Milk') or (Item.category = 'BEVERAGES'))  (cost=101.50 rows=102) (actual time=0
                                -> Table scan on Item  (cost=101.50 rows=1000) (actual time=0.042..0.377 rows=1000 loops=1)
                            -> Index lookup on li using item_id (item_id=Item.item_id)  (cost=0.40 rows=2) (actual time=0.003..0.004 rows=
                -> Single-row index lookup on l using PRIMARY (listing_id=`<subquery3>`.listing_id)  (cost=40.73 rows=1) (actual time=0.00
            -> Filter: (Recommendation.donor_id = l.listed_by)  (cost=0.79 rows=2) (actual time=0.001..0.001 rows=0 loops=105)
```

The overall cost of the query dropped from 486.07 (no index) to 309.85 (with index)

2) Indexing on Item.category only: We see a cost drop from 486.07 (no index) to 469.65

```
CREATE INDEX idx_category_only ON Item(category);

EXPLAIN ANALYZE SELECT DISTINCT l.listed_by,
                l.listing_id,
                l.location,
                l.status
FROM Listing l
LEFT JOIN (
    SELECT donor_id, score
```

Output | #CREATE INDEX idx_ca...ly ON Item(category); | #drop index idx_item_name_only on Item;

EXPLAIN

```
-> Table scan on <temporary>  (cost=469.65..480.27 rows=651) (actual time=1.509..1.526 rows=105 loops=1)
    -> Temporary table with deduplication  (cost=469.63..469.63 rows=651) (actual time=1.507..1.507 rows=105 loops=1)
        -> Nested loop left join  (cost=404.53 rows=651) (actual time=1.020..1.379 rows=105 loops=1)
            -> Nested loop inner join  (cost=338.86 rows=289) (actual time=1.011..1.225 rows=105 loops=1)
                -> Table scan on <subquery3>  (cost=231.60..237.69 rows=289) (actual time=1.000..1.014 rows=105 loops=1)
                    -> Materialize with deduplication  (cost=231.58..231.58 rows=289) (actual time=0.999..0.999 rows=105 loops=1)
                        -> Nested loop inner join  (cost=202.67 rows=289) (actual time=0.082..0.971 rows=109 loops=1)
                            -> Filter: ((Item.item_name = 'Milk') or (Item.category = 'BEVERAGES'))  (cost=101.50 rows=182) (actual time=0
                                -> Table scan on Item  (cost=101.50 rows=1000) (actual time=0.045..0.364 rows=1000 loops=1)
                            -> Index lookup on li using item_id (item_id=Item.item_id)  (cost=0.40 rows=2) (actual time=0.003..0.004 rows=
                -> Single-row index lookup on l using PRIMARY (listing_id=`<subquery3>`.listing_id)  (cost=72.36 rows=1) (actual time=0.00
            -> Filter: (Recommendation.donor_id = l.listed_by)  (cost=0.79 rows=2) (actual time=0.001..0.001 rows=0 loops=105)
```

3) Composite index (item_name, category) on item:

```
CREATE INDEX idx_item_name_category_only ON Item(item_name,category);

EXPLAIN ANALYZE SELECT DISTINCT l.listed_by,
                l.listing_id,
                l.location,
                l.status
FROM Listing l
LEFT JOIN (
    SELECT donor_id, score
    FROM Recommendation
    WHERE recipient_id = 1
```

Output | # drop index idx_category_only on Item; | Result 21

1 row | CSV

EXPLAIN

```
-> Table scan on <temporary>  (cost=309.85..316.90 rows=366) (actual time=1.649..1.663 rows=105 loops=1)
    -> Temporary table with deduplication  (cost=309.83..309.83 rows=366) (actual time=1.647..1.647 rows=105 loops=1)
        -> Nested loop left join  (cost=273.22 rows=366) (actual time=1.160..1.516 rows=105 loops=1)
            -> Nested loop inner join  (cost=236.05 rows=163) (actual time=1.149..1.400 rows=105 loops=1)
                -> Table scan on <subquery3>  (cost=174.67..179.16 rows=163) (actual time=1.135..1.149 rows=105 loops=1)
                    -> Materialize with deduplication  (cost=174.64..174.64 rows=163) (actual time=1.132..1.132 rows=105 loops=1)
                        -> Nested loop inner join  (cost=158.39 rows=163) (actual time=0.221..1.103 rows=109 loops=1)
                            -> Filter: ((Item.item_name = 'Milk') or (Item.category = 'BEVERAGES'))  (cost=101.50 rows=102) (actual time=0
                                -> Covering index scan on Item using idx_item_name_category_only  (cost=101.50 rows=1000) (actual time=0.1
                            -> Index lookup on li using item_id (item_id=Item.item_id)  (cost=0.40 rows=2) (actual time=0.003..0.004 rows=
                -> Single-row index lookup on l using PRIMARY (listing_id=`<subquery3>`.listing_id)  (cost=40.73 rows=1) (actual time=0.00
            -> Filter: (Recommendation.donor_id = l.listed_by)  (cost=0.79 rows=2) (actual time=0.001..0.001 rows=0 loops=105)
```

With the Composite Index, we observed some drop in the query cost from 486.07 to 309.85

Analysis :

In short, we can infer that having an index only on item_name provides the same results in cost compared to having a composite index of item_name with category. This can be because the cardinality of item_name is much more than category. Hence, more rows are filtered out while using an index on item_name.

AQ3: Indexing Design and Analysis

Cost of AQ3 with EXPLAIN ANALYZE before indexing:

```
EXPLAIN ANALYZE SELECT d.donor_id,
       u.name, u.location, u.email, u.phone,
       rev.avg_rating,
       rev.review_count
FROM Donor d
LEFT JOIN User u  1..n<->1:  ON d.user_id = u.user_id
LEFT JOIN (
   SELECT donor_id,
          Avg(rating) AS avg_rating,
          Count(*) AS review_count
   FROM Review r
   GROUP BY r.donor_id
) AS rev
ON d.donor_id = rev.donor_id
WHERE avg_rating IS NOT NULL AND review_count IS NOT NULL;
```

Services

Result 30 | Output

1 row

EXPLAIN

```
-> Nested loop left join  (cost=815.00 rows=1000) (actual time=2.957..4.971 rows=422 loops=1)
    -> Nested loop inner join  (cost=465.00 rows=1000) (actual time=2.943..3.843 rows=422 loops=1)
        -> Filter: (rev.donor_id is not null)  (cost=303.56..115.00 rows=1000) (actual time=2.920..3.130 rows=422 loops=1)
            -> Table scan on rev  (cost=303.76..318.75 rows=1000) (actual time=2.919..3.078 rows=422 loops=1)
                -> Materialize  (cost=303.75..303.75 rows=1000) (actual time=2.916..2.916 rows=422 loops=1)
                    -> Filter: (avg(r.rating) is not null)  (cost=203.75 rows=1000) (actual time=0.430..2.763 rows=422 loops=1)
                        -> Group aggregate: avg(r.rating), count(0)  (cost=203.75 rows=1000) (actual time=0.428..2.731 rows=422 loops=1)
                            -> Index scan on r using donor_id  (cost=103.75 rows=1000) (actual time=0.395..2.547 rows=1000 loops=1)
        -> Single-row index lookup on d using PRIMARY (donor_id=rev.donor_id)  (cost=0.25 rows=1) (actual time=0.001..0.002 rows=1 loops=422)
    -> Single-row index lookup on u using PRIMARY (user_id=d.user_id)  (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=422)
```

1. Creating an index on rating:

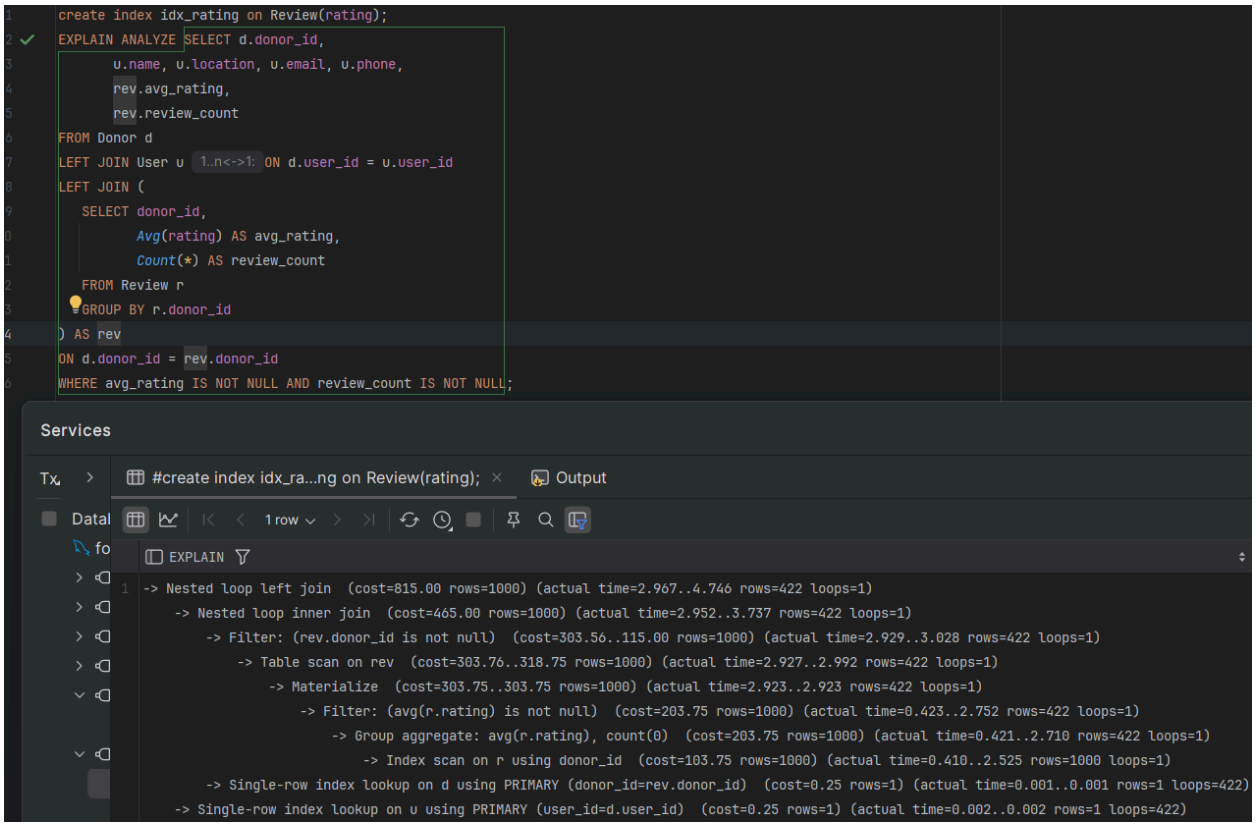

After creating an index on Rating, the overall cost of the query remains the same.

Analysis/Reason for no change in cost:
This result is likely due to the fact that the main filtering and join conditions in this query do not directly leverage the **review.rating** index. Specifically, the query retrieves aggregated data on **avg_rating** and **review_count** at the donor level, which involves grouping on **donor_id** rather than filtering by individual rating values. Consequently, the database engine processes the aggregation independently of the **review.rating** index, rendering it less effective for optimizing this query.

Note: The above query **only** has a **rating** field on which the indexing can be applied on. Other columns(**donor_id, user_id** etc) are already foreign or primary keys, hence we can't index on those. Indexing on other select columns like name, location also may not have any impact since they are not part of WHERE, GROUP BY clause.

AQ4: Indexing Design and Analysis

Cost of AQ4 with EXPLAIN ANALYZE before indexing:

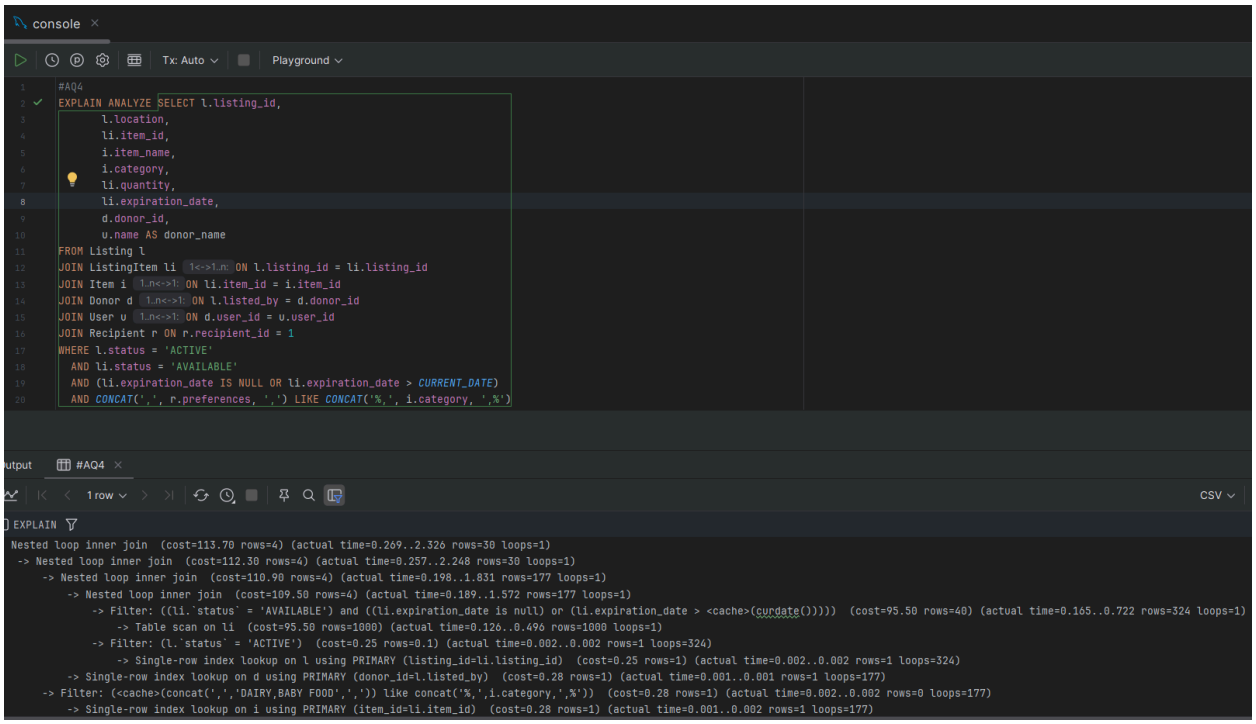

Let's Explore the index designs for AQ4:

1) Indexing on Listing.status Only:

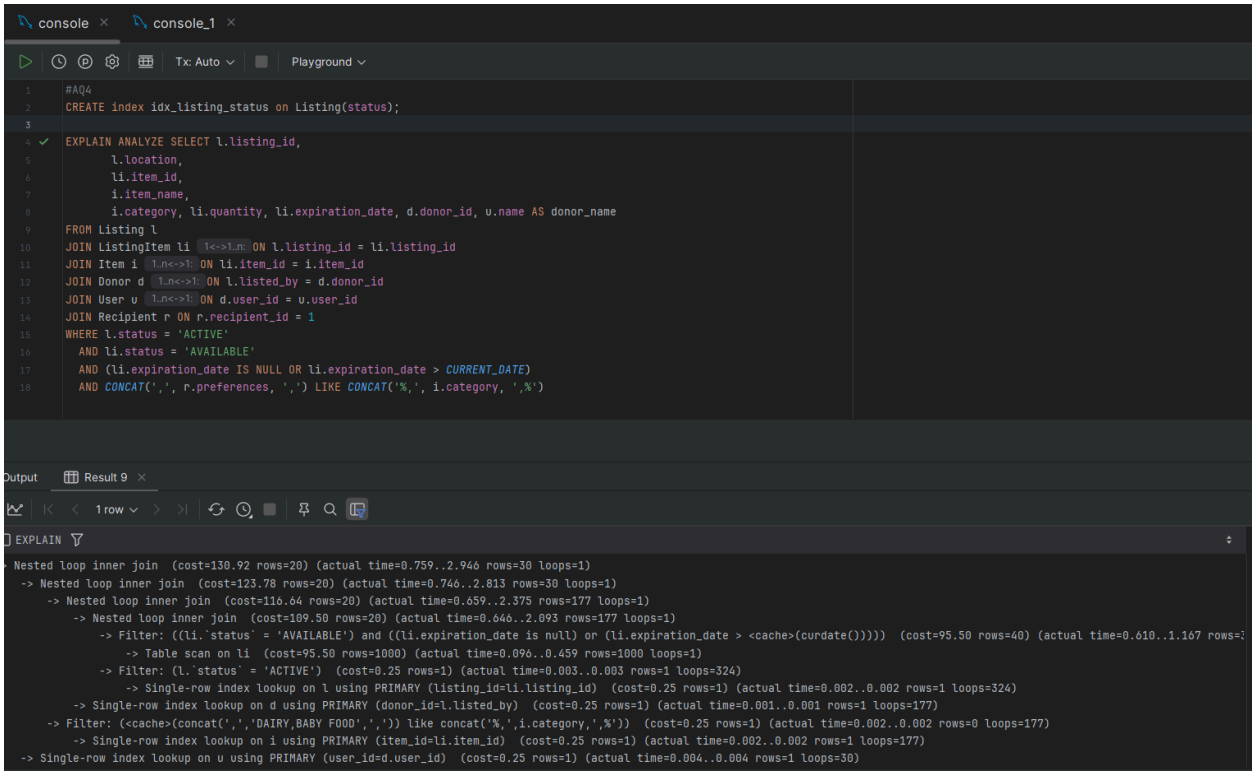

Creating an index on listing.status **increased** the query cost from 113.70 (no index) to 130.92. Adding an index on **Listing.status** increased the query cost. This is likely because the filtering

condition WHERE **l.status = 'ACTIVE'** in the query has low selectivity, meaning that many records in Listing meet this condition. Since the status is not highly selective, the database engine may still need to scan a large portion of Listing records, making the index less effective.

2) Indexing on ListingItem.status only:

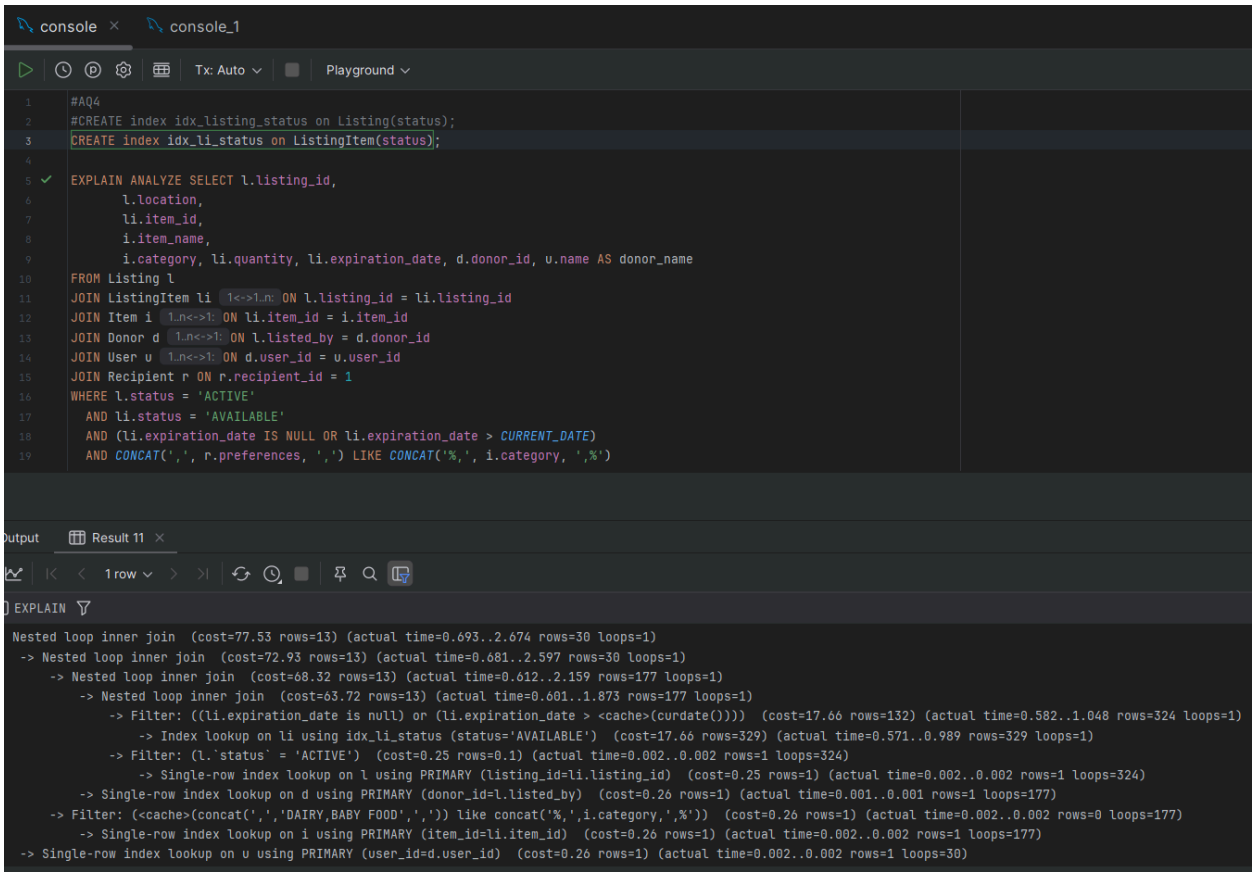

Creating an index on ListingItem.status **substantially dropped** the overall cost from 113.70 (no index) to 77.53. This means creating an index on **ListingItem.status** provided a significant performance benefit, as WHERE **li.status = 'AVAILABLE'** is likely a selective filter in our query. This index allows the database to quickly locate only the AVAILABLE items within ListingItem, reducing the number of rows that need to be joined with other tables.

3) Indexing on ListingItem.expiration_date:

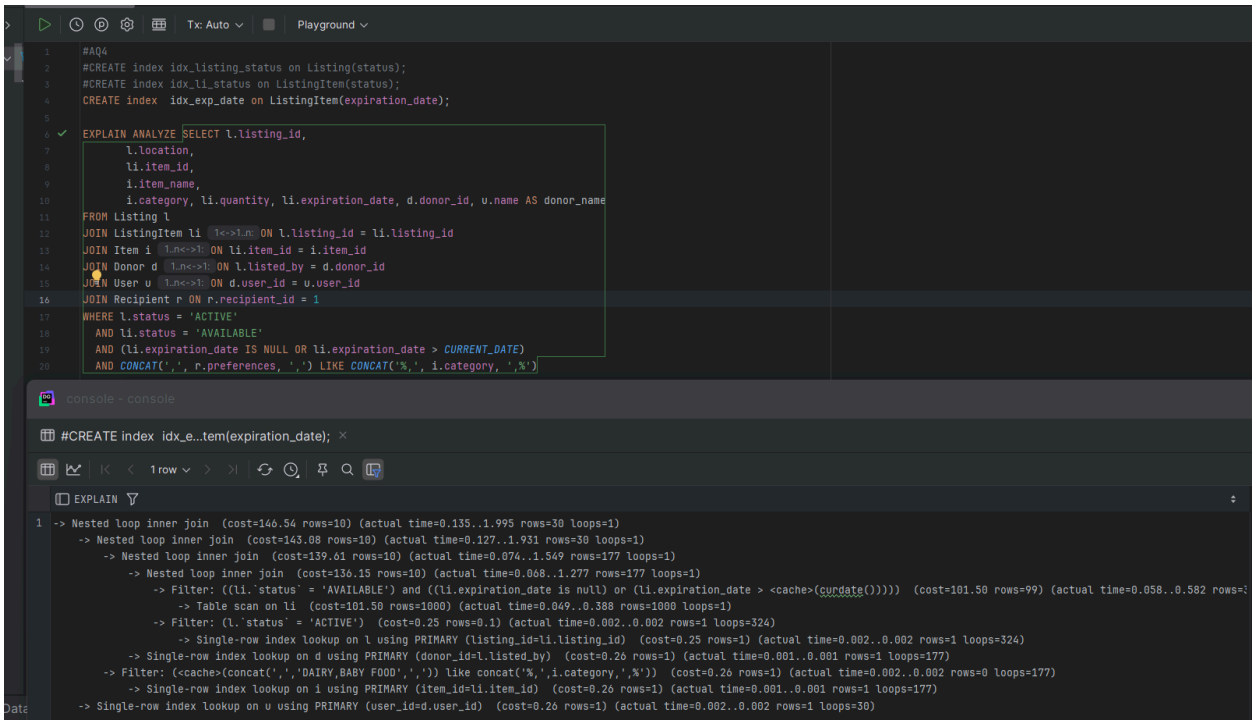

Indexing on ListingItem.expiration_date **increased** the overall cost from 113.70 to 146.54. This is likely because the condition (**li.expiration_date IS NULL OR li.expiration_date > CURRENT_DATE**) requires a complex filter, where the index isn't efficiently utilized due to the need to check for both **NULL** values and dates greater than the current date. The **OR** condition complicates index usage since it doesn't support range scanning efficiently.

Analysis:
In short, we can conclude that the index design having index on **`ListingItem.status`** works best compared to other index design approaches experimented.

Optimization Summary

Best index designs are selected based on the analysis given in the Indexing section.

Query	Best designed Index selected	EXPLAIN ANALYZE Overall Cost		Cost Optimization
		Before Indexing	After Indexing	
AQ1	ListingItem.status	5.44	5.44	No difference
AQ2	Item.item_name	486.07	309.85	Improved
AQ3	Review.rating	815.00	815.00	No difference
AQ4	ListingItem.status	113.70	77.53	Improved